

CT5132/CT5148 Week 12 Exercises

James McDermott

NUI Galway

Again this week, our exercises are extracted from the lecture slides/videos, and solutions are given below.

We won't look at the exercises for `ggplot`, as the Data Visualisation module is coming up soon where you'll see that in much more detail.

Exercises (dplyr joins)

- 1 Read the three data files `rentals.csv`, `movies.csv`, `customers.csv`, all in the `data/` directory, as tibbles.
- 2 Optional: get R to read the `Date` column correctly. Hint: https://readr.tidyverse.org/reference/parse_datetime.html
- 3 Using a `dplyr` join command, create a table showing the customer name and address for every rental.
- 4 Piping the result into another join command, recreate the full original table as shown under “Before Normalisation” above.
- 5 Notice the columns `Name.x` and `Name.y` which appear because there is a `Name` column in each of the `Movies` and `Customers` tables. Rename them.
- 6 Calculate the number of movies Frida watched of the Sci-fi genre.

Solutions (dplyr joins)

```
library(tidyverse)
```

```
## -- Attaching packages ----- tidyverse

## v ggplot2 3.1.0      v purrr  0.3.1
## v tibble  2.0.1      v dplyr  0.8.0.1
## v tidyr   0.8.3      v stringr 1.4.0
## v readr   1.3.1      v forcats 0.4.0

## Warning: package 'tibble' was built under R version 3.5.2
## Warning: package 'tidyr' was built under R version 3.5.2
## Warning: package 'purrr' was built under R version 3.5.2
## Warning: package 'dplyr' was built under R version 3.5.2
## Warning: package 'stringr' was built under R version 3.5.2
## Warning: package 'forcats' was built under R version 3.5.2
```

Exercises 1 and 2:

```
rentals <- read_csv("data/rentals.csv",  
                    col_types=cols(Date=col_date(  
                      format="%d-%b-%Y")))  
movies <- read_csv("data/movies.csv")
```

```
## Parsed with column specification:  
## cols(  
##   MovieID = col_double(),  
##   Name = col_character(),  
##   Genre = col_character()  
## )
```

```
customers <- read_csv("data/customers.csv")
```

```
## Parsed with column specification:  
## cols(  
##   CustomerID = col_double(),
```

Customer name and address for each rental

```
inner_join(rentals, customers, by="CustomerID")
```

```
## # A tibble: 5 x 5
```

##	Date	MovieID	CustomerID	Name	Address
##	<date>	<dbl>	<dbl>	<chr>	<chr>
## 1	2018-01-01	102	1	Bob	11, Haight St
## 2	2018-01-02	101	2	Frida	Oxford Circus
## 3	2018-01-02	102	3	Carrie	99, Fifth Ave
## 4	2018-01-05	103	1	Bob	11, Haight St
## 5	2018-01-05	104	2	Frida	Oxford Circus

Recreate original table

```
inner_join(rentals, customers, by="CustomerID") %>%  
  inner_join(movies, by="MovieID")
```

```
## # A tibble: 5 x 7
```

##	Date	MovieID	CustomerID	Name.x	Address	Name.y
##	<date>	<dbl>	<dbl>	<chr>	<chr>	<chr>
## 1	2018-01-01	102	1	Bob	11, Haight St	Amelie
## 2	2018-01-02	101	2	Frida	Oxford Circus	The Ma
## 3	2018-01-02	102	3	Carrie	99, Fifth Ave	Amelie
## 4	2018-01-05	103	1	Bob	11, Haight St	Skyfal
## 5	2018-01-05	104	2	Frida	Oxford Circus	Avenge

Rename columns

```
t = inner_join(rentals, customers, by="CustomerID") %>%  
  inner_join(movies, by="MovieID") %>%  
  rename(CustomerName=Name.x, MovieTitle=Name.y)  
t
```

```
## # A tibble: 5 x 7
```

##	Date	MovieID	CustomerID	CustomerName	Address	Mo
##	<date>	<dbl>	<dbl>	<chr>	<chr>	<c
## 1	2018-01-01	102	1	Bob	11, Haught~	An
## 2	2018-01-02	101	2	Frida	Oxford Cir~	Th
## 3	2018-01-02	102	3	Carrie	99, Fifth ~	An
## 4	2018-01-05	103	1	Bob	11, Haught~	Sk
## 5	2018-01-05	104	2	Frida	Oxford Cir~	Av

Filter and count

```
t %>% filter(CustomerName=="Frida", Genre=="Sci-fi") %>%  
  count()
```

```
## # A tibble: 1 x 1  
##       n  
##   <int>  
## 1     2
```

Filter and count

The following is a solution to the problem, but it requires the programmer to do all the work in their head. That's not scalable or flexible and it's error-prone, so don't do this.

```
# Frida is CustomerID 2  
# Movies 101 and 104 are Sci-fi  
rentals %>% filter(CustomerID == 2,  
                    MovieID %in% c(101, 104)) %>%  
  count()
```

```
## # A tibble: 1 x 1  
##       n  
##   <int>  
## 1     2
```

Exercises

- 1 In the `mpg` dataset (part of the tidyverse), calculate the mean and standard deviation of the highway fuel efficiency.
- 2 Using `group_by`, calculate the mean and standard deviation of the highway fuel efficiency per manufacturer.
- 3 Calculate the correlation between highway fuel efficiency and engine size.
- 4 What was the average highway fuel efficiency in 1999 and in 2008?
- 5 Carry out a two-sample independent t-test between highway fuel efficiency in 1999 and 2008 and interpret the result.
- 6 Carry out a regression on highway fuel efficiency by displacement.

Solution 1

```
library(tidyverse)
```

```
mean(mpg$hwy)
```

```
## [1] 23.44017
```

```
sd(mpg$hwy)
```

```
## [1] 5.954643
```

Solution 2

```
mpg %>% group_by(manufacturer) %>%  
  summarise(mean=mean(hwy), sd=sd(hwy))
```

```
## # A tibble: 15 x 3  
##   manufacturer mean    sd  
##   <chr>         <dbl> <dbl>  
## 1 audi          26.4  2.18  
## 2 chevrolet     21.9  5.11  
## 3 dodge         17.9  3.57  
## 4 ford          19.4  3.33  
## 5 honda         32.6  2.55  
## 6 hyundai       26.9  2.18  
## 7 jeep          17.6  3.25  
## 8 land rover    16.5  1.73  
## 9 lincoln       17     1  
## 10 mercury      18     1.15  
## 11 nissan        24.6  5.00
```

Solution 3

```
cor(mpg$hwy, mpg$displ)
```

```
## [1] -0.76602
```

Solution 4

```
mpg %>% group_by(year) %>%  
  summarise(mean=mean(hwy), sd=sd(hwy))
```

```
## # A tibble: 2 x 3  
##   year mean    sd  
##   <int> <dbl> <dbl>  
## 1  1999  23.4  6.08  
## 2  2008  23.5  5.85
```

Solution 5

```
mpg1999 <- mpg %>% filter(year == 1999)
mpg2008 <- mpg %>% filter(year == 2008)
t.test(mpg1999$hwy, mpg2008$hwy)
```

```
##
## Welch Two Sample t-test
##
## data: mpg1999$hwy and mpg2008$hwy
## t = -0.032864, df = 231.64, p-value = 0.9738
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
##  -1.562854  1.511572
## sample estimates:
## mean of x mean of y
##  23.42735  23.45299
```


Solution 6

```
res = lm(hwy ~ displ, data=mpg)
summary(res)
```

```
##
```

```
## Call:
```

```
## lm(formula = hwy ~ displ, data = mpg)
```

```
##
```

```
## Residuals:
```

```
##      Min       1Q   Median       3Q      Max
## -7.1039 -2.1646 -0.2242  2.0589 15.0105
```

```
##
```

```
## Coefficients:
```

```
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  35.6977     0.7204   49.55  <2e-16 ***
## displ       -3.5306     0.1945  -18.15  <2e-16 ***
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```