

Introduction to IR Models

Introduction to Information Retrieval Models

Information Retrieval vs. Data Retrieval

Data collections

Well structured collections of related items; items are usually atomic with a well-defined interpretation. Data retrieval involves the selection of a fixed set of data based on a well-defined query (e.g SQL, OQL).

Information collections

Information, on the other hand, is usually semi-structured or unstructured. Information retrieval (IR) involves the retrieval of documents of natural language which is typically not structured and may be semantically ambiguous.

Information Retrieval vs. Information Filtering

The main differences are:

- the nature of the information need
- the nature of the document set

Similarities

Other than these two differences, the same models are used. Documents and queries are represented using the same set of techniques and similar comparison algorithms are also used.

User Role

In traditional IR, the user role was pretty well-defined in that a user:

- formulated a query
- viewed the results
- potentially offered feedback
- possibly reformulated query and repeated steps

User Role

In more recent systems, with the increasing popularity of the hypertext paradigm, users usually intersperses browsing with the traditional querying.
Raises many new difficulties and challenges

Pre-processing

Document Pre-processing

Application of a set of well-known techniques to the documents and queries prior to any comparison. Includes, among others:

- stemming
- stop-word removal
- thesaurus construction

Stemming

Stemming refers to the reduction of words to a potentially common root.

Example

Computerisation, computing, computers could all be stemmed to common form *comput*.

Stemming involves the reduction of similar words to a common root form. Lovin's and Porter's algorithms are the most common.

Stop word removal

- This involves the removal of very frequent terms from documents.
- These terms add little to the semantics or meaning of the document.

Thesaurus construction

Thesauri used to try to identify synonyms within the documents. Manually or automatically created.

Representation

Representation and comparison technique depends on the information retrieval model chosen. The choice of feedback techniques is also dependent on the model chosen.

Models

IR Models

Retrieval models can be broadly categorised as:

- Boolean
 - Classical Boolean
 - Fuzzy Set approach
 - Extended Boolean
- Vector
 - Vector space approach
 - Latent Semantic Indexing
 - Neural Networks
- Probabilistic
 - Inference Network
 - Belief Network

IR model

Can view any IR model as comprising:

- D is the set of logical representations of the documents.
- Q is the the set of logical representations of the user information needs (queries).
- F is a framework for modelling these representations (D and Q) and the relationship between D and Q .
- R is a ranking function which defines an ordering among the documents with regard to any query q_i .

IR models

We have a set of index terms:

$$t_1 \dots t_n$$

A weight $w_{i,j}$ is assigned to each term t_i occurring in document d_j

Can view a document or query as a vector of weights:

$$\vec{d_j} = (w_1, w_2, w_3 \dots)$$

Boolean Model

Boolean Model

- Based on set theory and the Boolean algebra.
- A query is viewed as a Boolean expression.
- The model also assumes terms are present or absent, hence term weights $w_{i,j}$ are binary and discrete, i.e., $w_{i,j}$ is an element of $\{0, 1\}$
- Suffers from quite a few shortcomings:
 - people often have difficulty formulating expressions
 - documents are considered either relevant or irrelevant; no partial matching allowed

Example

$q = t1 \text{ AND } (t2 \text{ OR } (\text{NOT } t3))$

This can be mapped to what is termed disjunctive normal form, where we have a series of disjunctions (or logical ORs) of conjunctions.

$$q = 100 \vee 110 \vee 111$$

If a document 'satisfies' any of the components, the document is deemed relevant and returned.

Advantages

Clean formalism popular, widespread relatively simple

Disadvantages

- Not very good performance.
- Suffers badly from natural language effects of synonymy etc.
- No ranking of results.
- Harbours some difficulty in use.
- Terms in a documents are considered independent of each other.

Vector Space model

Vector space model

Attempts to improve upon the Boolean model by removing the limitation of binary weights for index terms.

Terms can have a non-binary weights in both queries and documents.

Hence we can represent documents and query as n-dimensional vectors.

$$\vec{d}_j = (w_{1,j}, w_{2,j} \dots w_{n,j})$$

$$\vec{q} = (w_{1,q}, w_{2,q} \dots w_{n,q})$$

Similarity

We can calculate the similarity between a document and a query by calculating the similarity between the vector representations of the document and query.

We can measure this similarity by measuring the cosine of the angle between the two vectors.

$$\vec{a} \cdot \vec{b} = |\vec{a}| |\vec{b}| \cos(\vec{a}, \vec{b})$$

$$\Rightarrow \cos(\vec{a}, \vec{b}) = \frac{\vec{a} \cdot \vec{b}}{|\vec{a}| |\vec{b}|}$$

We can therefore calculate similarity between document and query as:

$$\text{sim}(q, d) = \cos(\vec{q}, \vec{d}) = \frac{\vec{q} \cdot \vec{d}}{|\vec{q}| |\vec{d}|}$$

Considering term weights on query and document, we can calculate similarity between document and query as:

$$\text{sim}(q, d) = \frac{\sum_{i=1}^N (w_{i,q} \times w_{i,d})}{\sqrt{\sum_{i=1}^N (w_{i,q})^2} \times \sqrt{\sum_{i=1}^N (w_{i,d})^2}}$$

Weighting schemes

- We need a means to calculate the term weights in the document and query vector representations.
- A term's frequency within a document quantifies how well a term describes a document. The more frequent a term occurs in a document, the better it is at describing that document and vice-versa.
- This frequency is known as the term frequency or *tf* factor.

Weighting schemes

- If a term occurs frequently across all the documents, that term does little to distinguish one document from another. This factor is known as the inverse document frequency (*idf*-frequency).
- Traditionally, the most commonly used weighting schemes are known as *tf-idf* weighting schemes.

For all terms in a document, the weight assigned can be calculated as:

$$w_{i,j} = f_{i,j} \times \log\left(\frac{N}{N_i}\right)$$

where

- $f_{i,j}$ is the (possibly normalised) frequency of term t_i in document d_j
- N is the number of documents in the collection
- N_i is the number of documents that contain term t_i .

Advantages

Improved performance over the Boolean model due to weighting schemes Partial matching allowed which gives a natural ranking

Disadvantages

Terms are considered to be mutually independent